A DECLARATIVE SYSTEM FOR OPTIMIZING AI WORKLOADS

A PREPRINT

Chunwei Liu*, Matthew Russo*, Michael Cafarella, Lei Cao†, Peter Baille Chen, Zui Chen, Michael Franklin‡, Tim Kraska, Samuel Madden, Gerardo Vitagliano

MIT, †University of Arizona, ‡University of Chicago

chunwei@mit.edu, mdrusso@mit.edu, michjc@csail.mit.edu
caolei@arizona.edu, peterbc@mit.edu, chenz429@mit.edu, mjfranklin@uchicago.edu,
kraska@mit.edu, madden@csail.mit.edu, gerarvit@csail.mit.edu

Semantic Analytics Applications

A homebuyer wants to use online real estate (a)modern and attractive (b)within two miles of MIT listing data to find a place that is

A researcher in cancer research domain would like to

(a) download a dozen of cancer research papers

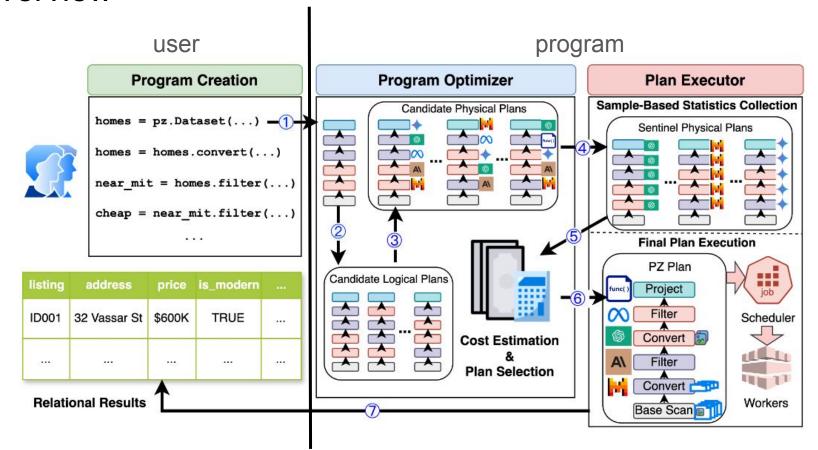
(b)identify the papers that contain patient experiment data

(c)integrate those data into a single table

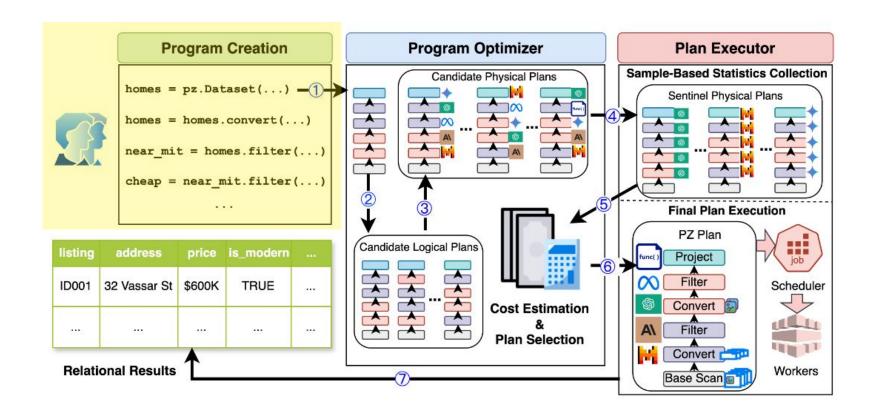
Semantic Analytics Applications

- Traditional data processing + Al-like semantic reasoning
- Data-intensive
- Can be decomposed into an execution tree of distinct operations
- May result in answers of varying quality

Overview



Program Creation



Custom schema in a relational view

```
class Email (pz.TextFile):

"""Represents an email, which can subclass a text file"""

sender = pz.StringField(desc="The email address of the sender", required=True)

subject = pz.StringField(desc="The subject of the email", required=True)

# define logical plan
emails = pz.Dataset(source="enron-emails", schema=Email) # invokes a convert operation
emails = emails.filter("The email is not quoting from a news article or an article ...")
emails = emails.filter("The email refers to a fraudulent scheme (i.e., \"Raptor\", ...")

# user specified policy
policy = pz.MinimizeCostAtFixedQuality(min_quality=0.8)

# execute plan
results = pz.Execute(emails, policy=policy)
```

Figure 3: The AI program written using PALIMPZEST for the Legal Discovery workload.

```
import palimpzest as pz

class Email(pz.TextFile):
    """Represents an email, which can subclass a text file"""
    sender = pz.StringField(desc="The email address of the sender", required=True)
    subject = pz.StringField(desc="The subject of the email", required=True)

# define logical plan
    emails = pz.Dataset(source="enron-emails", schema=Email) # invokes a convert operation
    emails = emails.filter("The email is not quoting from a news article or an article ...")
    emails = emails.filter("The email refers to a fraudulent scheme (i.e., \"Raptor\", ...")

# user specified policy
    policy = pz.MinimizeCostAtFixedQuality(min_quality=0.8)

# execute plan
```

results = pz.Execute (emails, policy=policy)

Figure 3: The AI program written using PALIMPZEST for the Legal Discovery workload.

Natural language descriptions

- For developers to specify correct program output
- For systems to find a high-quality implementation
- No need for prompt engineering from user side
- What if the given description contains error?
 - Skip this record processing
- What if it's ambiguous?
 - Explained in correctness section

Logical relational operators

Convert

- Transform a typed data object into a new object with specified schema
- The user does not need to specify how to implement a convert operation
- If LLM is needed.

operator	description
Project	π (rel., cols)
Select	σ (rel., predicate)
Convert	χ (rel., schema_a, schema_b)
Group By	Γ(rel., group_cond., agg.)
Limit	L(rel., limit)
Agg.	α(rel., agg_func)

(b) PALIMPZEST's full relational algebra. We extend the traditional relational algebra to include operators such as groupby which produce multiple relations.

```
import palimpzest as pz
   class Email (pz. TextFile) :
      """Represents an email, which can subclass a text file"""
     sender = pz.StringField(desc="The email address of the sender", required=True)
     subject = pz.StringField(desc="The subject of the email", required=True)
   # define logical plan
   emails = pz.Dataset(source="enron-emails", schema=Email) # invokes a convert operation
   emails = emails.filter("The email is not quoting from a news article or an article ...")
   emails = emails.filter("The email refers to a fraudulent scheme (i.e., \"Raptor\", ...")
   # user specified policy
13
14
   policy = pz.MinimizeCostAtFixedQuality(min_quality=0.8)
15
16
   # execute plan
   results = pz.Execute (emails, policy=policy)
```

Figure 3: The AI program written using PALIMPZEST for the Legal Discovery workload.

Correctness & Quality

- Correctness
 - What if the description is ambiguous?
 - Modify to be more precise
 - Split into more concrete steps
 - Future: validation example from user side

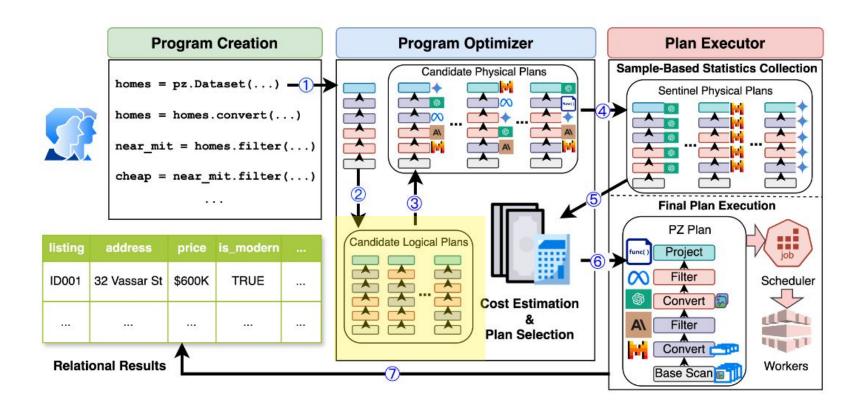
Quality

- How to assess the quality of an execution plan?
- How to find approaches to maximize assessed data quality?
- Compare with GPT-4 result as evaluation in experiment

```
import palimpzest as pz
   class Email (pz. TextFile) :
      """Represents an email, which can subclass a text file"""
     sender = pz.StringField(desc="The email address of the sender", required=True)
     subject = pz.StringField(desc="The subject of the email", required=True)
   # define logical plan
   emails = pz.Dataset(source="enron-emails", schema=Email) # invokes a convert operation
   emails = emails.filter("The email is not quoting from a news article or an article ...")
   emails = emails.filter("The email refers to a fraudulent scheme (i.e., \"Raptor\", ...")
12
   # user specified policy
13
   policy = pz.MinimizeCostAtFixedQuality(min_quality=0.8)
15
   # execute plan
   results = pz.Execute (emails, policy=policy)
```

Figure 3: The AI program written using PALIMPZEST for the Legal Discovery workload.

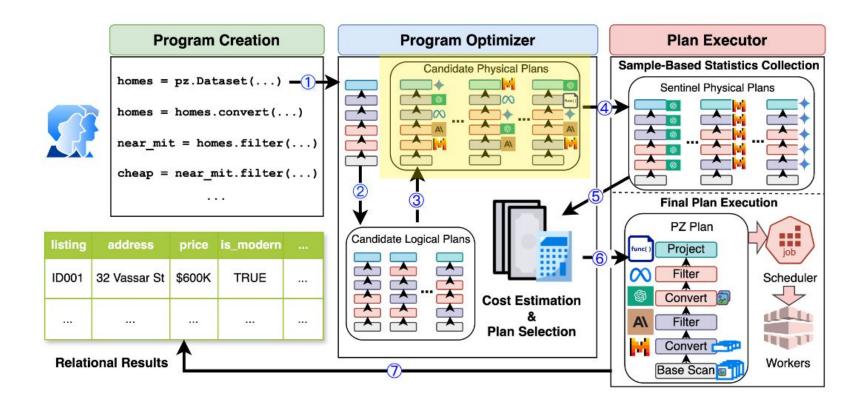
Program Optimizer



Logical Plan Optimization

- Filter reordering
 - Permutes the ordering of selection filters (3 steps -> 6 plans)
- Convert reordering
 - Move "convert" operations around the logical plan

Program Optimizer



Physical Plan Optimization Methods

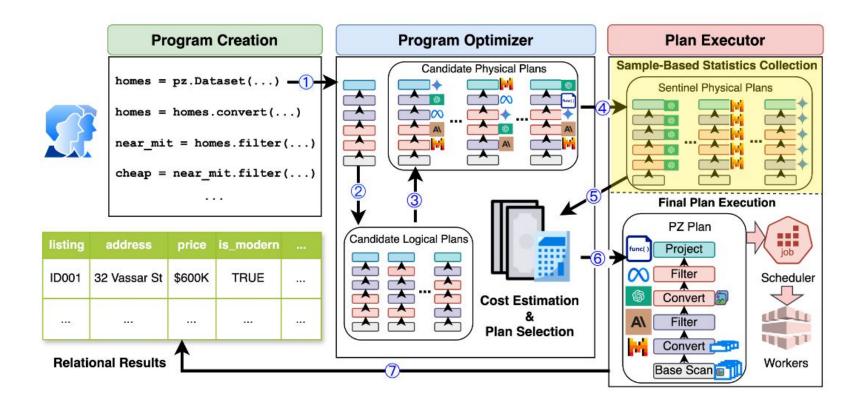
Implemented

- Model selection
- Code synthesis
- Multi-data prompt marshalling
- Input token reduction

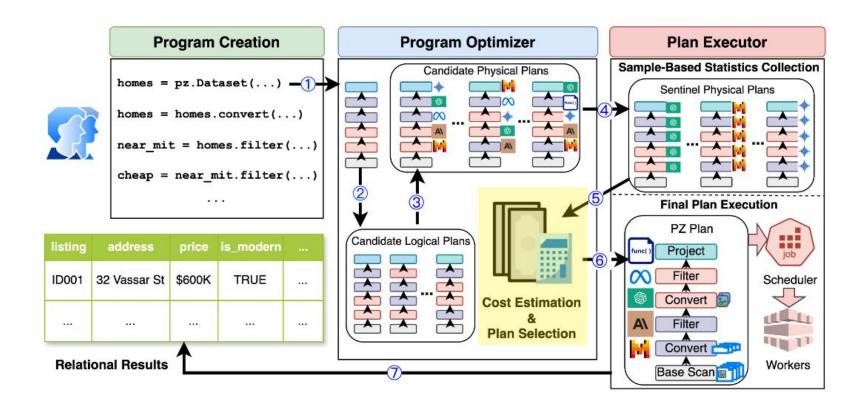
Future plan

- Output token reduction
- Model cascades
- Knowledge distillation
- Workload-aware execution management

Sample based Statistics Collection



Cost Estimation & Plan Selection

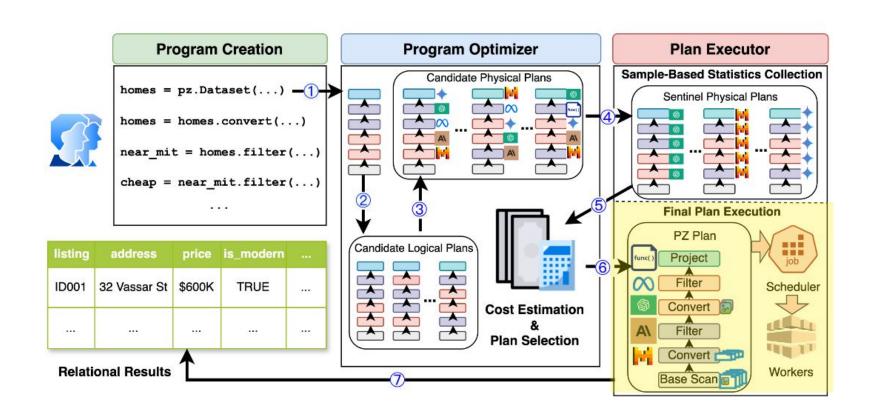


Cost Estimation & Plan Selection

```
Algorithm 1 Optimized Plan Selection Algorithm
Require: userCode, userPolicy
                                                                                                # Step (1)

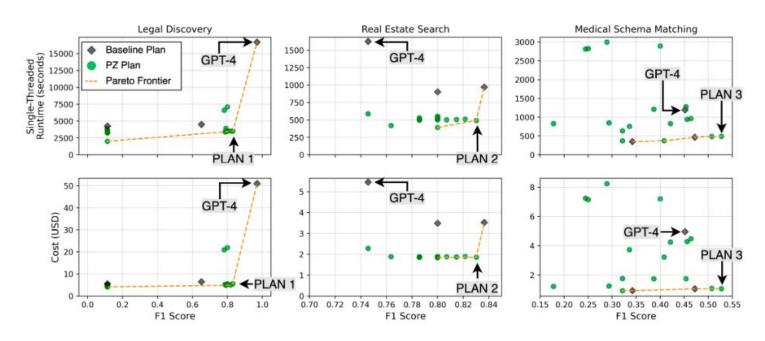
    logicalPlans = generateLogicalCandidates(userCode)

                                                                                                # Step ②
 2: sentinelPhysicalPlans = getPhysicalPlans(logicalPlans, sentinel = True)
                                                                                                # Step ③
 3:
 4: performanceStatisics = {}
 5: for 0...NUM_SAMPLES do
     input = getSampledInput()
 6:
     stats = runAndComputeStatistics(sentinelPhysicalPlans, input)
                                                                                                # Step 4
     performanceStatistics.update(stats)
                                                                                                 # Step (5)
9: end for
10:
11: physicalCandidates = getPhysicalPlans(logicalPlans, stats = performanceStatistics)
12: reducedCandidates = naiveElimination(physicalCandidates)
13: frontierCandidates = scoreAndEliminatePlans(reducedCandidates, performanceStatistics)
14:
                                                                                                 # Step 6
15: return chooseBestPlan(frontierCandidates, userPolicy)
```



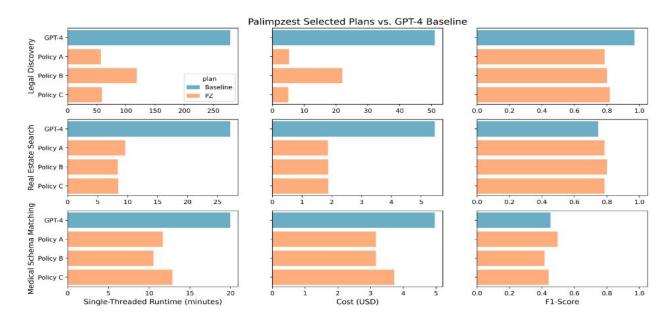
Experiments - 1

Palimpzest can generate plans which provide users with compelling performance trade-offs



Experiments - 2

 Palimpzest can identify plans that have better end-to-end runtime, cost and quality than a naive plan that uses the same state-of-the-art language model for each operation



Experiments - 3

 They ran Palimpzest with parallel implementations of the convert and filter operations to demonstrate the system's ability to achieve large runtime speedups

(a) Legal Discovery

Plan	Runtime (s)	Cost (\$)	F1
Single-Threaded Baseline (GPT-4)	16,712	51.0	0.97
Palimpzest	185 (1.1%)	5.60 (11.0%)	0.81 (83.5%)

(b) Real Estate Search

Plan	Runtime (s)	Cost (\$)	F1
Single-Threaded Baseline (GPT-4)	1,626	5.46	0.75
Palimpzest	80.9 (5.0%)	1.86 (34.1%)	0.80 (107%)

(c) Medical Schema Matching

Plan	Runtime (s)	Cost (\$)	F1
Single-Threaded Baseline (GPT-4)	1,195	4.96	0.45
Palimpzest	215 (18.0%)	3.36 (67.7%)	0.46 (102%)